

SurvMeth 640 Assignment 3

Cheng, Chia Wen

2023-03-13

Setup

```
library(mlbench)
library(rpart)
library(partykit)
library(caret)
```

Data

In this notebook, we use the Boston Housing data set (again). “This dataset contains information collected by the U.S Census Service concerning housing in the area of Boston Mass. It was obtained from the StatLib archive (<http://lib.stat.cmu.edu/datasets/boston>), and has been used extensively throughout the literature to benchmark algorithms.”

Source: <https://www.cs.toronto.edu/~delve/data/boston/bostonDetail.html>

```
data(BostonHousing2)
head(BostonHousing2)
```

```
##      town tract      lon      lat medv cmedv      crim zn indus chas   nox
## 1   Nahant  2011 -70.9550 42.2550 24.0  24.0 0.00632 18  2.31    0 0.538
## 2 Swampscott 2021 -70.9500 42.2875 21.6  21.6 0.02731  0  7.07    0 0.469
## 3 Swampscott 2022 -70.9360 42.2830 34.7  34.7 0.02729  0  7.07    0 0.469
## 4 Marblehead 2031 -70.9280 42.2930 33.4  33.4 0.03237  0  2.18    0 0.458
## 5 Marblehead 2032 -70.9220 42.2980 36.2  36.2 0.06905  0  2.18    0 0.458
## 6 Marblehead 2033 -70.9165 42.3040 28.7  28.7 0.02985  0  2.18    0 0.458
##      rm age    dis rad tax ptratio      b lstat
## 1 6.575 65.2 4.0900  1 296    15.3 396.90  4.98
## 2 6.421 78.9 4.9671  2 242    17.8 396.90  9.14
## 3 7.185 61.1 4.9671  2 242    17.8 392.83  4.03
## 4 6.998 45.8 6.0622  3 222    18.7 394.63  2.94
## 5 7.147 54.2 6.0622  3 222    18.7 396.90  5.33
## 6 6.430 58.7 6.0622  3 222    18.7 394.12  5.21
```

```
names(BostonHousing2)
```

```
## [1] "town"      "tract"     "lon"       "lat"       "medv"      "cmedv"     "crim"
## [8] "zn"        "indus"     "chas"      "nox"       "rm"        "age"       "dis"
## [15] "rad"       "tax"       "ptratio"   "b"         "lstat"
```

Drop some variables that are not needed.

```
BostonHousing2$town <- NULL
BostonHousing2$tract <- NULL
BostonHousing2$cmedv <- NULL
```

Split the data into a train and test set.

```
set.seed(9384)

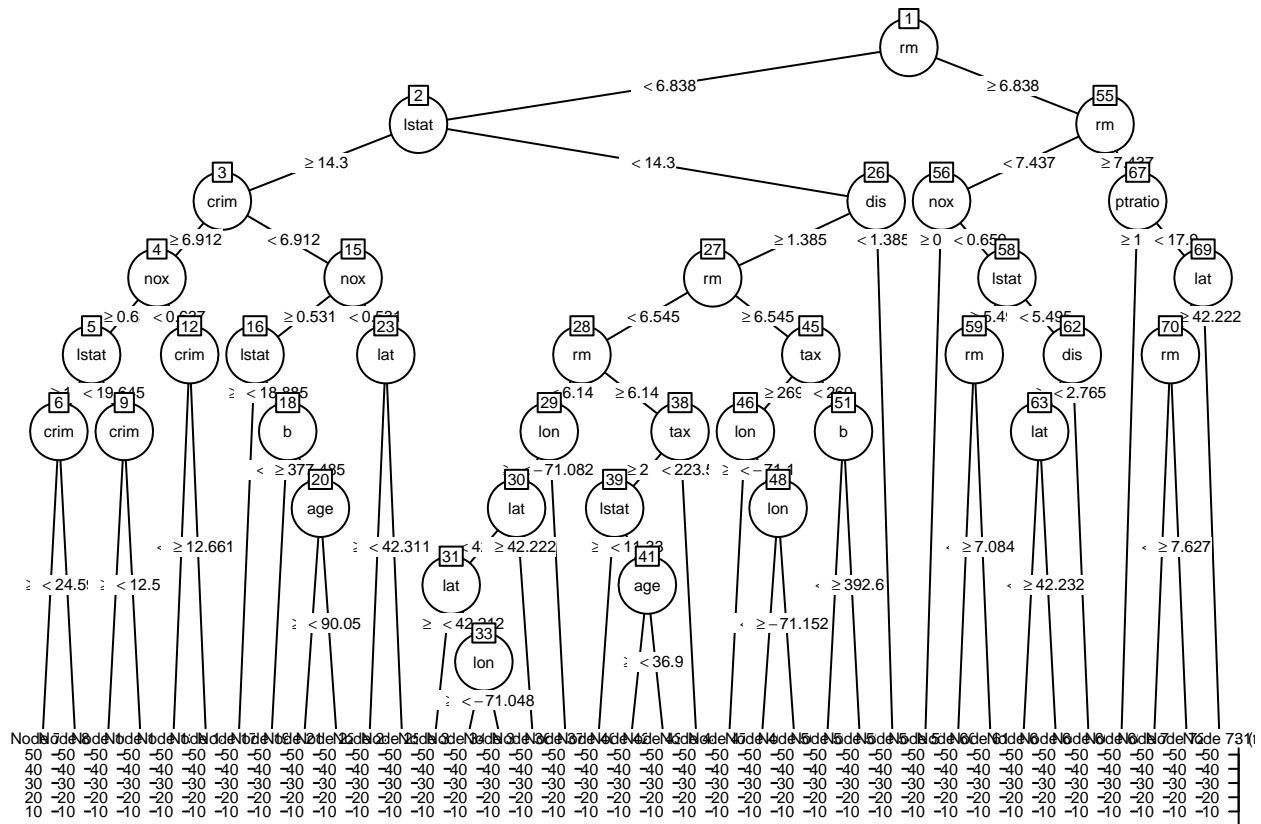
train <- sample(1:nrow(BostonHousing2), 0.8*nrow(BostonHousing2))
boston_train <- BostonHousing2[train,]
boston_test <- BostonHousing2[-train,]
```

CART

As before, we are interested in training a model for predicting the median home values (`medv`), using all features available. However, in this notebook we want to use CART as the prediction method. Make sure to grow a large tree, since we want to prune it back later.

```
# build a larger tree
tree2 <- rpart(medv ~ ., data = boston_train, method = "anova",
               control = rpart.control(minsplit = 10, # minimal obs in a node
                                       minbucket = 3, # minimal obs in any terminal node
                                       cp = 0.001, # min improvement through splitting
                                       maxdepth = 30 # maximum tree depth
               ))

party_tree2 <- as.party(tree2)
plot(party_tree2, gp = gpar(fontsize = 6))
```




This large tree is likely to overfit and will not generalize well to new data. Therefore, we use `printcp()` and `plotcp()` that help us to determine the best subtree.

```
printcp(tree2)
```

```
##
## Regression tree:
## rpart(formula = medv ~ ., data = boston_train, method = "anova",
##       control = rpart.control(minsplit = 10, minbucket = 3, cp = 0.001,
##                               maxdepth = 30))
##
## Variables actually used in tree construction:
## [1] age      b      crim    dis      lat      lon      lstat    nox      ptratio
## [10] rm      tax
##
## Root node error: 35535/404 = 87.959
##
## n= 404
##
##          CP nsplit rel error  xerror   xstd
## 1  0.4565952      0  1.000000  1.00141  0.090237
## 2  0.1596255      1  0.543405  0.66770  0.067684
## 3  0.1040341      2  0.383779  0.43916  0.054033
## 4  0.0615025      3  0.279745  0.38854  0.047334
## 5  0.0290994      4  0.218243  0.29586  0.044003
## 6  0.0259984      5  0.189143  0.27048  0.043780
```

## 7	0.0252640	6	0.163145	0.26759	0.043610
## 8	0.0069852	7	0.137881	0.24340	0.038564
## 9	0.0068438	8	0.130896	0.22216	0.037322
## 10	0.0063681	11	0.110364	0.22275	0.037401
## 11	0.0056780	12	0.103996	0.21917	0.037047
## 12	0.0055366	13	0.098318	0.21803	0.036753
## 13	0.0051020	14	0.092781	0.21845	0.036646
## 14	0.0048107	15	0.087679	0.21748	0.036644
## 15	0.0044513	16	0.082868	0.21411	0.036619
## 16	0.0040860	17	0.078417	0.21196	0.036880
## 17	0.0030686	18	0.074331	0.20729	0.036806
## 18	0.0030198	19	0.071263	0.21055	0.035551
## 19	0.0025376	20	0.068243	0.21114	0.036283
## 20	0.0018645	21	0.065705	0.20622	0.035827
## 21	0.0018179	22	0.063841	0.20773	0.035903
## 22	0.0016243	23	0.062023	0.20725	0.035896
## 23	0.0015804	24	0.060398	0.20828	0.035891
## 24	0.0014355	25	0.058818	0.20760	0.035863
## 25	0.0013252	26	0.057383	0.20801	0.035793
## 26	0.0013069	27	0.056057	0.20879	0.035884
## 27	0.0012211	28	0.054750	0.20926	0.035890
## 28	0.0011483	29	0.053529	0.20908	0.035896
## 29	0.0011469	30	0.052381	0.21017	0.036527
## 30	0.0011407	31	0.051234	0.21017	0.036527
## 31	0.0010317	32	0.050093	0.21037	0.036534
## 32	0.0010166	35	0.046998	0.21123	0.036584
## 33	0.0010000	36	0.045982	0.21108	0.036578

```
plotcp(tree2)
```



Briefly summarize the output of these two functions. What is plotted/printed here?

We prune the tree to avoid overfitting of the data. With the complexity parameter of 0.0018645, we get the least cross validated error of 0.24890. This is when 21 splits are applied to the decision tree. The dotted line in the plot points out where 'xerror' equals to 0.3. It shows that trees with more than 7 splits (not including 7) have 'xerrors' lower than 0.3 and the change is minor, but a smaller 'cp' does not guarantee a lower 'xerror.'

end

On this basis, we are interested in picking the cp value that is associated with the smallest CV error. Store this value in an object.

```
minx <- which.min(tree2$cptable[, "xerror"])
mincp <- tree2$cptable[minx, "CP"]
mincp # 0.001864483
```

```
## [1] 0.001864483
```

In addition, we could also pick the best subtree based on the 1-SE rule. Again, store the corresponding cp value in an object.

```
minx <- which.min(tree2$cptable[, "xerror"])
minxse <- tree2$cptable[minx, "xerror"] + tree2$cptable[minx, "xstd"]
minse <- which(tree2$cptable[1:minx, "xerror"] < minxse)
mincp2 <- tree2$cptable[minse[1], "CP"]
mincp2 # 0.006843847
```

```
## [1] 0.006843847
```

Now we can get the best subtree with the `prune()` function. First based on the smallest CV error...

```
p_tree_small_CV <- prune(tree2, cp = mincp)
p_tree_small_CV
```

```
## n= 404
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 404 35535.34000 22.686140
##    2) rm< 6.8375 330 13274.68000 19.685150
##      4) lstat>=14.3 139 2657.36200 14.825180
##        8) crim>=6.91188 59 763.50750 11.649150
##          16) nox>=0.6365 47 440.21320 10.659570
##            32) lstat>=19.645 35 199.06570 9.542857 *
##            33) lstat< 19.645 12 70.19667 13.916670 *
##          17) nox< 0.6365 12 97.00250 15.525000 *
##          9) crim< 6.91188 80 859.79550 17.167500
##            18) nox>=0.531 64 535.68480 16.373440
##              36) lstat>=18.885 21 138.89810 14.123810 *
##              37) lstat< 18.885 43 238.60650 17.472090 *
##            19) nox< 0.531 16 122.33940 20.343750 *
##          5) lstat< 14.3 191 4944.96800 23.221990
##            10) dis>=1.38485 188 2759.45500 22.794680
##              20) rm< 6.5445 151 1339.51900 21.697350
##                40) rm< 6.1405 76 478.05740 20.423680
##                  80) lon>=-71.08215 59 325.12030 19.838980 *
##                  81) lon< -71.08215 17 62.76235 22.452940 *
##                41) rm>=6.1405 75 613.23920 22.988000
##                  82) tax>=223.5 71 380.33320 22.657750
##                    164) lstat>=11.33 14 31.11429 20.157140 *
##                    165) lstat< 11.33 57 240.17510 23.271930 *
##                  83) tax< 223.5 4 87.71000 28.850000 *
##                21) rm>=6.5445 37 496.07300 27.272970
##                  42) tax>=269 26 230.71120 25.773080
##                    84) lon>=-71.1 15 59.25333 24.033330 *
##                    85) lon< -71.1 11 64.14727 28.145450 *
##                  43) tax< 269 11 68.61636 30.818180 *
##                11) dis< 1.38485 3 0.00000 50.000000 *
##          3) rm>=6.8375 74 6035.39900 36.068920
##            6) rm< 7.437 50 1947.16100 31.172000
##              12) nox>=0.659 3 27.92000 14.400000 *
##              13) nox< 0.659 47 1021.47500 32.242550
```

```
##          26) lstat>=5.495 26    368.39880 30.265380
##          52) rm< 7.0835 14     72.85429 27.357140 *
##          53) rm>=7.0835 12     38.98917 33.658330 *
##          27) lstat< 5.495 21    425.59810 34.690480
##          54) dis>=2.7648 18     102.14940 33.294440 *
##          55) dis< 2.7648 3      77.88667 43.066670 *
##          7) rm>=7.437 24      391.34960 46.270830
##          14) ptratio>=17.9 4     46.94750 40.125000 *
##          15) ptratio< 17.9 20    163.10000 47.500000 *
```

...and now based on the 1-SE rule.

```
p_tree_one_SE <- prune(tree2, cp = mincp2)
p_tree_one_SE
```

```
## n= 404
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 404 35535.3400 22.68614
##    2) rm< 6.8375 330 13274.6800 19.68515
##      4) lstat>=14.3 139 2657.3620 14.82518
##        8) crim>=6.91188 59 763.5075 11.64915 *
##        9) crim< 6.91188 80 859.7955 17.16750 *
##      5) lstat< 14.3 191 4944.9680 23.22199
##        10) dis>=1.38485 188 2759.4550 22.79468
##          20) rm< 6.5445 151 1339.5190 21.69735
##            40) rm< 6.1405 76 478.0574 20.42368 *
##            41) rm>=6.1405 75 613.2392 22.98800 *
##          21) rm>=6.5445 37 496.0730 27.27297 *
##        11) dis< 1.38485 3 0.0000 50.00000 *
##    3) rm>=6.8375 74 6035.3990 36.06892
##      6) rm< 7.437 50 1947.1610 31.17200
##        12) nox>=0.659 3 27.9200 14.40000 *
##        13) nox< 0.659 47 1021.4750 32.24255 *
##      7) rm>=7.437 24 391.3496 46.27083 *
```

Now, plot the smaller tree.

```
tree1 <- rpart(medv ~ ., data = boston_train)
tree1
```

```
## n= 404
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 404 35535.3400 22.68614
##    2) rm< 6.8375 330 13274.6800 19.68515
##      4) lstat>=14.3 139 2657.3620 14.82518
##        8) crim>=6.91188 59 763.5075 11.64915 *
##        9) crim< 6.91188 80 859.7955 17.16750 *
```

```
##      5) lstat< 14.3 191 4944.9680 23.22199
##      10) rm< 6.5445 153 2920.6570 22.06732
##      20) lstat>=9.66 82 502.1888 20.58049 *
##      21) lstat< 9.66 71 2027.8330 23.78451
##      42) indus< 14.48 64 586.9511 22.88281 *
##      43) indus>=14.48 7 913.0943 32.02857 *
##      11) rm>=6.5445 38 998.9982 27.87105 *
##      3) rm>=6.8375 74 6035.3990 36.06892
##      6) rm< 7.437 50 1947.1610 31.17200
##      12) lstat>=9.65 8 409.2750 23.17500 *
##      13) lstat< 9.65 42 928.8190 32.69524 *
##      7) rm>=7.437 24 391.3496 46.27083 *
```

```
summary(tree1)
```

```
## Call:
## rpart(formula = medv ~ ., data = boston_train)
##      n= 404
##
##              CP nsplit rel error      xerror      xstd
## 1 0.45659520      0 1.0000000 1.0058112 0.09072378
## 2 0.15962553      1 0.5434048 0.6355163 0.06498087
## 3 0.10403412      2 0.3837793 0.4084787 0.04991205
## 4 0.02909945      3 0.2797452 0.3166916 0.04422542
## 5 0.02885333      4 0.2506457 0.2927256 0.04457137
## 6 0.01713975      5 0.2217924 0.2735420 0.04712084
## 7 0.01292266      6 0.2046526 0.2565365 0.04033433
## 8 0.01000000      8 0.1788073 0.2490509 0.04046544
##
## Variable importance
##      rm      lstat      indus      nox      dis      age      tax ptratio      crim      zn
##      32      22      9      7      6      5      5      5      3      2
##      rad      lon      lat
##      2      1      1
##
## Node number 1: 404 observations,      complexity param=0.4565952
##      mean=22.68614, MSE=87.95877
##      left son=2 (330 obs) right son=3 (74 obs)
##      Primary splits:
##      rm      < 6.8375      to the left, improve=0.4565952, (0 missing)
##      lstat      < 9.725      to the right, improve=0.4521823, (0 missing)
##      lon      < -71.0685      to the right, improve=0.2747556, (0 missing)
##      indus      < 7.225      to the right, improve=0.2716115, (0 missing)
##      ptratio      < 18.75      to the right, improve=0.2605731, (0 missing)
##      Surrogate splits:
##      lstat      < 4.83      to the right, agree=0.876, adj=0.324, (0 split)
##      ptratio      < 14.55      to the right, agree=0.847, adj=0.162, (0 split)
##      indus      < 3.985      to the right, agree=0.842, adj=0.135, (0 split)
##      zn      < 87.5      to the left, agree=0.829, adj=0.068, (0 split)
##      nox      < 0.4045      to the right, agree=0.822, adj=0.027, (0 split)
##
## Node number 2: 330 observations,      complexity param=0.1596255
##      mean=19.68515, MSE=40.22629
##      left son=4 (139 obs) right son=5 (191 obs)
```



```

## Primary splits:
##   lstat < 14.3      to the right, improve=0.4273059, (0 missing)
##   nox   < 0.6695   to the right, improve=0.3007320, (0 missing)
##   crim  < 9.311465 to the right, improve=0.2632838, (0 missing)
##   dis   < 2.37495  to the left,  improve=0.2323209, (0 missing)
##   ptratio < 19.9    to the right, improve=0.2276603, (0 missing)
## Surrogate splits:
##   nox < 0.5765     to the right, agree=0.800, adj=0.525, (0 split)
##   age < 83.6       to the right, agree=0.800, adj=0.525, (0 split)
##   indus < 16.57    to the right, agree=0.788, adj=0.496, (0 split)
##   dis < 2.37495   to the left,  agree=0.782, adj=0.482, (0 split)
##   tax < 434.5      to the right, agree=0.773, adj=0.460, (0 split)
##
## Node number 3: 74 observations,    complexity param=0.1040341
##   mean=36.06892, MSE=81.55944
##   left son=6 (50 obs) right son=7 (24 obs)
## Primary splits:
##   rm < 7.437       to the left,  improve=0.6125342, (0 missing)
##   lstat < 4.68      to the right, improve=0.3985944, (0 missing)
##   ptratio < 19.15   to the right, improve=0.2160262, (0 missing)
##   lon < -71.04625   to the right, improve=0.2078671, (0 missing)
##   rad < 16          to the right, improve=0.1549302, (0 missing)
## Surrogate splits:
##   lstat < 3.99      to the right, agree=0.824, adj=0.458, (0 split)
##   lon < -71.15755   to the right, agree=0.716, adj=0.125, (0 split)
##   zn < 81.25        to the left,  agree=0.703, adj=0.083, (0 split)
##   indus < 1.23      to the right, agree=0.703, adj=0.083, (0 split)
##   ptratio < 14.75   to the right, agree=0.703, adj=0.083, (0 split)
##
## Node number 4: 139 observations,    complexity param=0.02909945
##   mean=14.82518, MSE=19.11771
##   left son=8 (59 obs) right son=9 (80 obs)
## Primary splits:
##   crim < 6.91188    to the right, improve=0.3891299, (0 missing)
##   lstat < 19.605     to the right, improve=0.3345170, (0 missing)
##   dis < 2.0037       to the left,  improve=0.3337001, (0 missing)
##   nox < 0.657        to the right, improve=0.3321810, (0 missing)
##   tax < 567.5        to the right, improve=0.2824666, (0 missing)
## Surrogate splits:
##   rad < 16          to the right, agree=0.871, adj=0.695, (0 split)
##   tax < 567.5        to the right, agree=0.856, adj=0.661, (0 split)
##   lat < 42.212       to the left,  agree=0.755, adj=0.424, (0 split)
##   nox < 0.657        to the right, agree=0.755, adj=0.424, (0 split)
##   dis < 2.202        to the left,  agree=0.734, adj=0.373, (0 split)
##
## Node number 5: 191 observations,    complexity param=0.02885333
##   mean=23.22199, MSE=25.88988
##   left son=10 (153 obs) right son=11 (38 obs)
## Primary splits:
##   rm < 6.5445        to the left,  improve=0.20734470, (0 missing)
##   lstat < 9.54         to the right, improve=0.19906760, (0 missing)
##   dis < 1.6141        to the right, improve=0.10956960, (0 missing)
##   lon < -71.03785     to the right, improve=0.09574916, (0 missing)
##   chas splits as LR, improve=0.08694740, (0 missing)

```

```

## Surrogate splits:
##   lstat < 5.055      to the right, agree=0.869, adj=0.342, (0 split)
##   crim < 0.017895   to the right, agree=0.827, adj=0.132, (0 split)
##   zn < 87.5         to the left,  agree=0.812, adj=0.053, (0 split)
##   lon < -70.8315    to the left,  agree=0.806, adj=0.026, (0 split)
##   dis < 10.648      to the left,  agree=0.806, adj=0.026, (0 split)
##
## Node number 6: 50 observations,      complexity param=0.01713975
##   mean=31.172, MSE=38.94322
##   left son=12 (8 obs) right son=13 (42 obs)
##   Primary splits:
##     lstat < 9.65      to the right, improve=0.3127974, (0 missing)
##     ptratio < 18.95   to the right, improve=0.2101861, (0 missing)
##     rad < 16          to the right, improve=0.1555746, (0 missing)
##     crim < 2.935235   to the right, improve=0.1555746, (0 missing)
##     tax < 534.5       to the right, improve=0.1555746, (0 missing)
##   Surrogate splits:
##     crim < 7.393425   to the right, agree=0.92, adj=0.500, (0 split)
##     nox < 0.659       to the right, agree=0.90, adj=0.375, (0 split)
##     rad < 16          to the right, agree=0.90, adj=0.375, (0 split)
##     tax < 534.5       to the right, agree=0.90, adj=0.375, (0 split)
##     indus < 15.015    to the right, agree=0.88, adj=0.250, (0 split)
##
## Node number 7: 24 observations
##   mean=46.27083, MSE=16.30623
##
## Node number 8: 59 observations
##   mean=11.64915, MSE=12.9408
##
## Node number 9: 80 observations
##   mean=17.1675, MSE=10.74744
##
## Node number 10: 153 observations,      complexity param=0.01292266
##   mean=22.06732, MSE=19.08926
##   left son=20 (82 obs) right son=21 (71 obs)
##   Primary splits:
##     lstat < 9.66      to the right, improve=0.13374900, (0 missing)
##     crim < 7.24712    to the left,  improve=0.11433340, (0 missing)
##     lon < -71.03785   to the right, improve=0.08881250, (0 missing)
##     rm < 6.1405       to the left,  improve=0.08419838, (0 missing)
##     dis < 1.68515     to the right, improve=0.08213046, (0 missing)
##   Surrogate splits:
##     nox < 0.519       to the right, agree=0.771, adj=0.507, (0 split)
##     age < 48.1        to the right, agree=0.739, adj=0.437, (0 split)
##     dis < 4.48025     to the left,  agree=0.739, adj=0.437, (0 split)
##     crim < 0.098325   to the right, agree=0.725, adj=0.408, (0 split)
##     indus < 7.625     to the right, agree=0.719, adj=0.394, (0 split)
##
## Node number 11: 38 observations
##   mean=27.87105, MSE=26.28943
##
## Node number 12: 8 observations
##   mean=23.175, MSE=51.15937
##

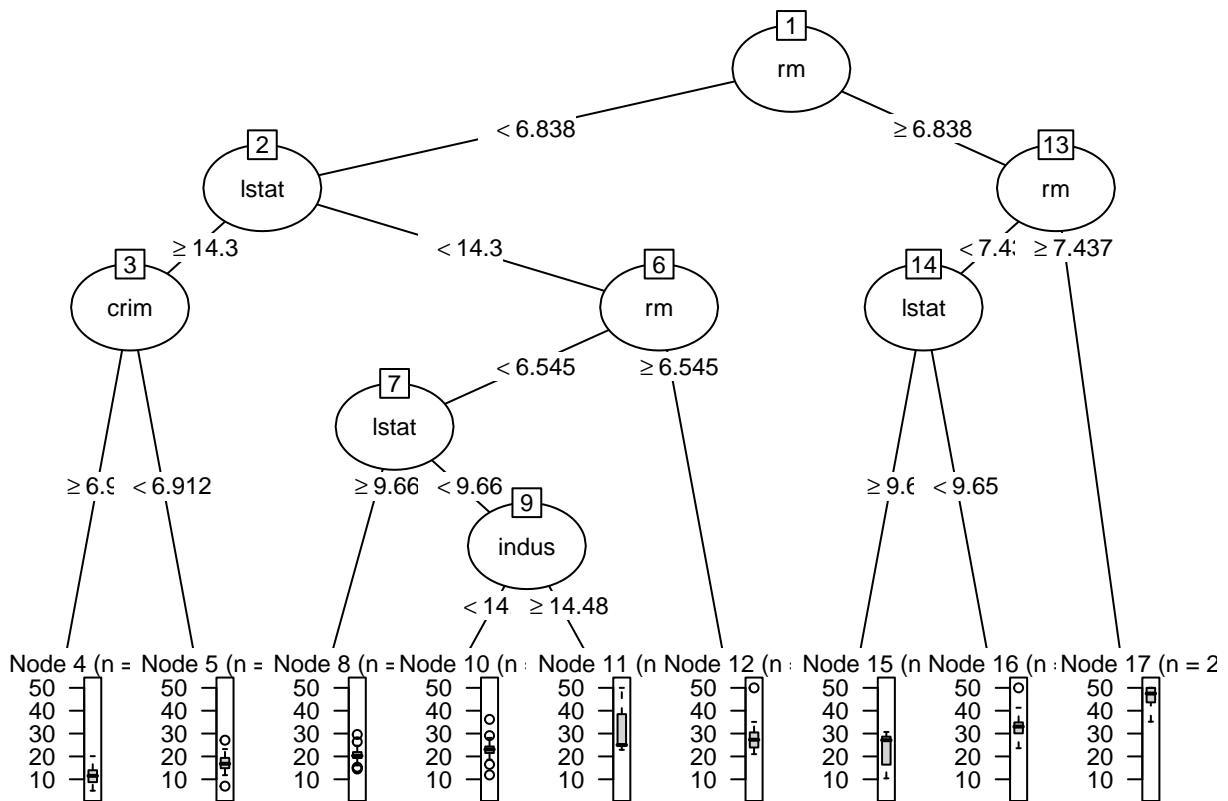
```

```

## Node number 13: 42 observations
##   mean=32.69524, MSE=22.11474
##
## Node number 20: 82 observations
##   mean=20.58049, MSE=6.124253
##
## Node number 21: 71 observations,   complexity param=0.01292266
##   mean=23.78451, MSE=28.56103
##   left son=42 (64 obs) right son=43 (7 obs)
##   Primary splits:
##       indus < 14.48      to the left,  improve=0.2602717, (0 missing)
##       age  < 81.8       to the left,  improve=0.2364868, (0 missing)
##       crim < 0.484795   to the left,  improve=0.2220958, (0 missing)
##       dis  < 2.6221     to the right, improve=0.2080417, (0 missing)
##       rad  < 6.5        to the left,  improve=0.1718898, (0 missing)
##   Surrogate splits:
##       crim < 1.163695   to the left,  agree=0.972, adj=0.714, (0 split)
##       nox  < 0.589      to the left,  agree=0.972, adj=0.714, (0 split)
##       age  < 86.05      to the left,  agree=0.972, adj=0.714, (0 split)
##       dis  < 2.04295    to the right, agree=0.958, adj=0.571, (0 split)
##       rad  < 16         to the left,  agree=0.944, adj=0.429, (0 split)
##
## Node number 42: 64 observations
##   mean=22.88281, MSE=9.171111
##
## Node number 43: 7 observations
##   mean=32.02857, MSE=130.442

party_tree1 <- as.party(tree1)
plot(party_tree1, gp = gpar(fontsize = 9))

```



Prediction

Finally, we can use the pruned trees to predict the outcome in the holdout (test) set.

```
y_tree_small_CV <- predict(p_tree_small_CV, newdata = boston_test, type = "vector")
y_tree_one_SE <- predict(p_tree_one_SE, newdata = boston_test, type = "vector")
```

Use at least two performance measures to evaluate the prediction performance of both trees.

```
# first performance measures
MSE_tree_small_CV <- mean((y_tree_small_CV-boston_test$medv)^2)
MSE_tree_small_CV ## 17.0959
```

```
## [1] 17.0959
```

```
MSE_tree_one_SE <- mean((y_tree_one_SE-boston_test$medv)^2)
MSE_tree_one_SE ## 21.35381
```

```
## [1] 21.35381
```

```
# second performance measures
eval_results <- function(true, predicted, df) {
  SSE <- sum((predicted - true)^2)
```

```

SST <- sum((true - mean(true))^2)
R_squared <- 1 - SSE / SST
RMSE = sqrt(SSE/nrow(df))

# Model performance metrics
data.frame(
  RMSE = RMSE,
  Rsquared = R_squared
)
}

eval_results(boston_test$medv, y_tree_small_CV, boston_test) ## RMSE: 4.134719; R-squared: 0.7555617

##          RMSE  Rsquared
## 1 4.134719 0.7555617

eval_results(boston_test$medv, y_tree_one_SE, boston_test) ## RMSE: 4.621019; R-squared: 0.6946818

##          RMSE  Rsquared
## 1 4.621019 0.6946818

```

Which tree would you recommend for prediction purposes?

For prediction purposes, I would recommend the tree with the smallest cross-validation value because it has a lower MSE of 17.0959 comparing to 21.35381 from the tree with 1-SE away from the smallest CV. Smaller RMSE indicates that the forecasted values of medv is closer to the actual medv within the testing data. The greater R-squared further proves that the tree with the smallest CV is better than the other one.

end